

# Welcome to the Machine

## Prerequisites

Yinbo Wang

School of Statistics, RUC

May 9, 2025

## Learning by doing

### Math

- **Linear Algebra**
- **Matrix Calculus**
- **Probability Theory**
- **Mathematical Statistics**

### Coding

- **Python**
- **Numpy**
- **PyTorch**

# 0-0 NumPy ('numpy')

## Numerical Python

### What is it?

- Fundamental package for scientific computing in Python.
- Provides a high-performance **multidimensional array object** and tools for working with these arrays.

### Core Object: 'ndarray'

- Fast, memory-efficient multidimensional array for *homogeneous* data.
- Dimensions:
  - 1D: Vector
  - 2D: Matrix
  - 3D+: Tensor

# 0-1 NumPy: Key Features & Why Use It

## Key Features

- **Vectorized Operations:** Element-wise ops on arrays without Python loops (**much faster**).
- **Broadcasting:** Operations on arrays of different shapes.
- **Rich Functionality:** Math functions, linear algebra, random numbers.
- **Interoperability:** Base for Pandas, SciPy, Scikit-learn, PyTorch.

## Why Use It?

- **Performance:** Significantly faster than pure Python.
- **Convenience:** Concise syntax for numerical operations.
- **Foundation:** Bedrock of scientific Python.

# 0-2 NumPy: Simple Example

## Basic Operations

```
import numpy as np

# Create a NumPy array from a Python list
a = np.array([1, 2, 3, 4, 5])
b = np.array([6, 7, 8, 9, 10])

# Element-wise addition (vectorized)
c = a + b
print(f"a: {a}")
print(f"b: {b}")
print(f"a + b: {c}")

# Scalar multiplication
d = a * 2
print(f"a * 2: {d}")

# Create a 2D array (matrix)
matrix = np.array([[1, 2, 3], [4, 5, 6]])
print(f"Matrix:\n{matrix}")
print(f"Matrix shape: {matrix.shape}")
```

# 0-3 PyTorch ('torch')

## Machine Learning Framework

### What is it?

- Open-source machine learning framework (Meta AI).
- Known for flexibility, ease of use, Python integration.
- Widely used for deep learning research and production.

### Core Object: 'Tensor'

- Multi-dimensional arrays, similar to NumPy's 'ndarray'.
- **Key Difference:** Can be moved to **GPUs for parallel computation**.
- Basis of **Automatic Differentiation (Autograd)**.

## 0-4 PyTorch: Key Features

- **GPU Acceleration:** Seamlessly run computations on GPUs.
- **Automatic Differentiation ('torch.autograd'):**
  - Automatically computes gradients.
  - Backbone of training neural networks via backpropagation.
  - Builds a "computational graph" dynamically.
- **Neural Network Module ('torch.nn'):** Pre-defined layers, loss functions.
- **Optimization Algorithms ('torch.optim'):** SGD, Adam, etc.
- **Utilities:** Data loading, distributed training.
- **Dynamic Computational Graphs:** Graph built "on the fly", easier debugging, flexible architectures.

# 0-5 PyTorch: Why Use It for ML/DL?

## Advantages

- **Python-first:** Natural for Python developers.
- **Flexibility & Control:** Balance of high-level abstractions and low-level control.
- **Strong Research Community:** Rapid adoption of new ideas.
- **Ease of Debugging:** Dynamic graphs aid inspection.



# 0-6 PyTorch: Simple Example

## Basic Operations

```
import torch
a = torch.tensor([1, 2, 3, 4, 5], dtype=torch.float32)
b = torch.tensor([6, 7, 8, 9, 10], dtype=torch.float32)

# Element-wise addition
c = a + b
print(f"a + b: {c}")

# Scalar multiplication
d = a * 2
print(f"a * 2: {d}")

# Create a 2D tensor (matrix)
matrix = torch.tensor([[1, 2, 3], [4, 5, 6]], dtype=torch.float32)
print(f"Matrix:\n{matrix}")
print(f"Matrix shape: {matrix.shape}")

# Automatic differentiation example
x = torch.tensor(2.0, requires_grad=True)
y = x**2 + 3*x + 1
y.backward() # Computes gradients dy/dx
print(f"x: {x}")
print(f"y = x^2 + 3x + 1: {y}")
print(f"dy/dx at x=2: {x.grad}") # Expected: 2*x + 3 = 2*2 + 3 = 7
```

## 0-7 NumPy vs. PyTorch Tensors

- PyTorch tensors can be easily converted to NumPy arrays and vice-versa.
- `tensor.numpy()`: PyTorch Tensor  $\rightarrow$  NumPy array.
- `torch.from_numpy(ndarray)`: NumPy array  $\rightarrow$  PyTorch Tensor.

### Interoperability

```
numpy_arr = np.array([10, 20, 30])
pytorch_tensor = torch.from_numpy(numpy_arr)
print(f"From NumPy to PyTorch: {pytorch_tensor}")

new_numpy_arr = pytorch_tensor.numpy()
print(f"From PyTorch to NumPy: {new_numpy_arr}")

# If a tensor is on GPU, you need to move it to CPU first
# if torch.cuda.is_available():
#     gpu_tensor = torch.tensor([1,2,3], device="cuda")
#     cpu_tensor = gpu_tensor.cpu() # Move to CPU
#     numpy_arr_from_gpu = cpu_tensor.numpy()
```

# 1-0 Linear Algebra

## Overview

Linear algebra is the branch of mathematics concerning vector spaces and linear mappings between them. In ML, it's fundamental for **representing data, defining transformations, and solving systems of equations**.

## Why it's important for ML/DL?

- **Data Representation:** Datasets as matrices, multimedia as tensors.
- **Model Parameters:** Weights and biases as matrices/vectors.
- **Transformations:** Core operations in NNs (matrix multiplications).
- **Dimensionality Reduction:** E.g., PCA.
- **Optimization:** Solving linear systems.

# 1-1 Scalars, Vectors, Matrices, Tensors

- **Scalar:** A single number (e.g., 5, 3.14).
- **Vector:** An ordered array of numbers.
  - In ML: Represents a single data point or feature vector.
- **Matrix:** A 2D array of numbers.
  - In ML: Represents a dataset or model parameters (NN layer weights).
- **Tensor:** Generalization to N dimensions.
  - 0D: Scalar, 1D: Vector, 2D: Matrix.
  - 3D: RGB image (height, width, channels).
  - 4D: Batch of RGB images (batch, height, width, channels).

## NumPy/PyTorch Implementation

Both libraries provide powerful tools for these:

- NumPy: `np.array()`, `.shape`, `.ndim`
- PyTorch: `torch.tensor()`, `.shape`, `.ndim`, `dtype`

# 1-2 Dot Product (Inner Product)

- For vectors  $\mathbf{a}, \mathbf{b}$  of length  $n$ :  $\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i$ . Result is a scalar.
- Geometric:  $\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \cdot \|\mathbf{b}\| \cos(\theta)$ .
  - Measures **similarity/alignment**. Orthogonal if  $\mathbf{a} \cdot \mathbf{b} = 0$ .

## NumPy/PyTorch Implementation

- NumPy: `np.dot(a,b)`, `a @ b`, `np.sum(a*b)`
- PyTorch: `torch.dot(a,b)`, `torch.matmul(a,b)`, `a @ b`

# 1-3: Matrix Multiplication

- $\mathbf{A} (m \times n), \mathbf{B} (n \times p) \implies \mathbf{C} = \mathbf{AB}$  is  $m \times p$ .
- $C_{ij} = \sum_{k=1}^n A_{ik}B_{kj}$  (dot product of  $i$ -th row of  $\mathbf{A}$  and  $j$ -th col of  $\mathbf{B}$ ).
- **Order matters:**  $\mathbf{AB} \neq \mathbf{BA}$  generally.
- Fundamental in NNs: `output = activation(weights @ inputs + bias)`.

## NumPy/PyTorch Implementation

- NumPy: `np.matmul(A,B), A @ B`
- PyTorch: `torch.matmul(A,B), A @ B, torch.mm(A,B)` (for 2D only)

# 1-4 Transpose

- $\mathbf{A}^T$ : Flips matrix over main diagonal. Rows become columns.
- If  $\mathbf{A}$  is  $m \times n$ ,  $\mathbf{A}^T$  is  $n \times m$ .  $(A^T)_{ij} = A_{ji}$ .
- Property:  $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$ .

## NumPy/PyTorch Implementation

- NumPy: `A.T`, `np.transpose(A)`
- PyTorch: `A.T`, `torch.transpose(A, 0, 1)`

# 1-5 Inverse & Pseudo-inverse

- **Inverse ( $\mathbf{A}^{-1}$ ):** For square  $\mathbf{A}$ , if it exists.
  - $\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$  (identity).
  - Only exists if  $\mathbf{A}$  is non-singular (determinant  $\neq 0$ ).
  - Solves  $\mathbf{A}\mathbf{x} = \mathbf{b} \implies \mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$ .
- **Pseudo-inverse ( $\mathbf{A}^+$ ):** Generalization for non-square or singular matrices.
  - Finds "best fit" (least squares) solution to  $\mathbf{A}\mathbf{x} = \mathbf{b}$ .
  - Used in Linear Regression (Normal Eq:  $\mathbf{w} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$ ).

## NumPy/PyTorch Implementation

- NumPy: `numpy.linalg.inv()`, `numpy.linalg.pinv()`
- PyTorch: `torch.linalg.inv()`, `torch.linalg.pinv()`



# 1-6 Determinant

- Scalar value from a square matrix.
- Geometrically: Signed area (2D) or volume (3D) of parallelogram/parallelepiped formed by column/row vectors.
- If determinant is 0, matrix is singular (no inverse, transformation collapses space).

## NumPy/PyTorch Implementation

- NumPy: `numpy.linalg.det()`
- PyTorch: `torch.linalg.det()`

# 1-7 Eigenvalues & Eigenvectors

- For square  $\mathbf{A}$ , eigenvector  $\mathbf{v}$  and eigenvalue  $\lambda$ :  $\mathbf{A}\mathbf{v} = \lambda\mathbf{v}$ .
- Vector  $\mathbf{v}$ 's direction is unchanged by  $\mathbf{A}$ , only scaled by  $\lambda$ .
- Eigenvectors: **Principal axes of transformation by  $\mathbf{A}$ .**
- Eigenvalues: **Scaling factor along these axes.**
- Crucial for PCA, matrix powers, stability analysis.

## NumPy/PyTorch Implementation

- NumPy: `numpy.linalg.eig()`
- PyTorch: `torch.linalg.eig()` (complex), `torch.linalg.eigh()` (real symmetric)

# 1-8: Singular Value Decomposition (SVD)

- Factorization of *any*  $m \times n$  matrix  $\mathbf{A}$  into  $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ .
  - $\mathbf{U}$ :  $m \times m$  orthogonal (left singular vectors).
  - $\mathbf{\Sigma}$ :  $m \times n$  diagonal (singular values  $\sigma_i \geq 0$ ).
  - $\mathbf{V}^T$ :  $n \times n$  orthogonal (rows are right singular vectors).
- Singular values  $\approx$  "strengths" of principal components.
- **Applications:** PCA, dimensionality reduction, matrix approximation, pseudo-inverse computation.
- Reconstructing  $\mathbf{A}$  from  $U, S, Vh$  is common: `U @ Sigma_matrix @ Vh`.

## NumPy/PyTorch Implementation

- NumPy: `numpy.linalg.svd()`
- PyTorch: `torch.linalg.svd()`

# 2-0 Matrix Calculus

## Overview

Extends calculus concepts (derivatives, gradients) to functions involving matrices and vectors. Essential for optimizing ML models.

## Why it's important for ML/DL?

- **Optimization:** Most ML models are trained by minimizing a loss function. Gradient Descent requires computing gradients of loss w.r.t. model parameters (matrices/vectors).
- **Backpropagation:** Algorithm for training NNs, an application of chain rule from matrix calculus.

## 2-1 Gradient

- For a **scalar-valued function**  $f(\mathbf{x})$  of a vector  $\mathbf{x} = [x_1, \dots, x_n]^T$ .
- Gradient  $\nabla_{\mathbf{x}} f(\mathbf{x})$  is a vector of partial derivatives:

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

- Points in direction of **steepest ascent**.  $-\nabla_{\mathbf{x}} f(\mathbf{x})$  points in direction of **steepest descent**.

## 2-2 Jacobian

- For a **vector-valued function**  $\mathbf{f}(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}^m$ .
- Jacobian matrix  $\mathbf{J}$  is an  $m \times n$  matrix of all first-order partial derivatives:  $J_{ij} = \frac{\partial f_i}{\partial x_j}$ .

$$\mathbf{J} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} = \begin{bmatrix} (\nabla_{\mathbf{x}} f_1(\mathbf{x}))^T \\ \vdots \\ (\nabla_{\mathbf{x}} f_m(\mathbf{x}))^T \end{bmatrix}$$

- If  $m = 1$  (scalar function), Jacobian is transpose of gradient vector.

## 2-3 Hessian

- For a **scalar-valued function**  $f(\mathbf{x})$  of a vector  $\mathbf{x}$ .
- Hessian matrix  $\mathbf{H}$  is an  $n \times n$  matrix of second-order partial derivatives:  $H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$ .

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

- Describes **local curvature** of the function. Used in second-order optimization (e.g., Newton's method).

## 2-4 Chain Rule (for vectors/matrices)

- Scalar case: If  $y = f(u)$  and  $u = g(\mathbf{x})$ , then  $\frac{\partial y}{\partial x_i} = \frac{df}{du} \frac{\partial g}{\partial x_i}$ .
- Vector case: If  $\mathbf{z} = f(\mathbf{y})$  and  $\mathbf{y} = g(\mathbf{x})$ , then the Jacobian of  $\mathbf{z}$  w.r.t.  $\mathbf{x}$  is:

$$\mathbf{J}_{\mathbf{x}}(\mathbf{z}) = \mathbf{J}_{\mathbf{y}}(\mathbf{z})\mathbf{J}_{\mathbf{x}}(\mathbf{y})$$

- Fundamental principle behind backpropagation in neural networks.



## 2-5 PyTorch Implementation (Autograd)

### Core Idea: 'autograd'

PyTorch's 'autograd' package automatically computes gradients.

- Tensors can track operations if `requires_grad=True`.
- A computation graph is built dynamically.
- `loss.backward()` computes gradients of 'loss' w.r.t. all parameters with `requires_grad=True`.
- Gradients are stored in the `.grad` attribute of tensors.

### Scalar Output Example (Conceptual)

```
x = torch.tensor([1.,2.,3.])
w = torch.tensor([0.1,0.2,0.3], requires_grad=True)

y = (w @ x)**2 # Forward pass, builds graph
y.backward()   # Backward pass, computes dy/dw

print(w.grad)  # Access gradient dy/dw
```

## 2-5 PyTorch Implementation (Autograd)

- For full Jacobians/Hessians:  
    `torch.autograd.functional.jacobian`,  
    `torch.autograd.functional.hessian`.
- `'backward()'` computes gradients essential for optimization via gradient descent.

### Vector Output Example (Conceptual)

```
x = torch.tensor([1.,2.,3.])
w = torch.tensor([0.1,0.2,0.3], requires_grad=True)

# For Jacobian
# Define the function whose Jacobian is sought
def compute_model_output(weights_arg):
    # 'x' is captured from the outer scope where it's defined.
    # 'weights_arg' will be the 'w' tensor when 'jacobian' is called.
    return torch.stack([weights_arg @ x, (weights_arg**2).sum()])

jacobian = torch.autograd.functional.jacobian(compute_model_output, w)
print(jacobian) # Access Jacobian
```

# 3-0 Probability Theory

## Quantifying Uncertainty and Randomness

### Concept

The mathematical framework for **quantifying uncertainty and randomness**. In its modern, axiomatic formulation, probability is a type of measure.

### Why it's important for ML/DL?

- **Modeling Uncertainty:** Data is noisy, processes are stochastic.
- **Generative Models:** Learn probability distributions (GANs, VAEs).
- **Loss Functions:** Derived from probabilistic principles (e.g., cross-entropy).
- **Bayesian Methods:** Probabilistic approach to inference.

## 3-1 Probability Space $(\Omega, \mathcal{F}, P)$

Models a random process with three components:

- **Sample Space  $(\Omega)$ :** Set of *all possible outcomes*.
  - Coin flip:  $\Omega = \{H, T\}$ . Dice roll:  $\Omega = \{1, 2, 3, 4, 5, 6\}$ .
- **Event Space  $(\mathcal{F}$  or Sigma-Algebra):** Collection of subsets of  $\Omega$  (events) to which we assign probabilities. Must satisfy:
  - 1  $\Omega \in \mathcal{F}$
  - 2 If  $A \in \mathcal{F}$ , then  $A^c \in \mathcal{F}$  (closed under complement)
  - 3 If  $A_1, A_2, \dots \in \mathcal{F}$ , then  $\bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$  (closed under countable unions)
- **Probability Measure  $(P)$ :** Function  $P : \mathcal{F} \rightarrow [0, 1]$  satisfying Kolmogorov Axioms:
  - 1 Non-negativity:  $P(A) \geq 0$  for any event  $A$ .
  - 2 Normalization:  $P(\Omega) = 1$ .
  - 3 Countable Additivity: For disjoint events  $A_i$ ,  $P(\bigcup A_i) = \sum P(A_i)$ .

## 3-2 Random Variables (RV)

- **Intuition:** Assigns a numerical value to each outcome in  $\Omega$ .
- **Formal Definition:** A function  $X : \Omega \rightarrow \mathbb{R}$  such that for every Borel set  $B \subset \mathbb{R}$  (any "reasonable" subset), its pre-image  $X^{-1}(B) = \{\omega \in \Omega : X(\omega) \in B\}$  is an event in  $\mathcal{F}$ .
  - This is **measurability**. Ensures we can ask  $P(X \in B)$ .
  - Sufficient to check for  $B = (-\infty, x]$ , i.e.,  $\{\omega : X(\omega) \leq x\} \in \mathcal{F}$ . This allows defining CDF.
- **Types:**
  - **Discrete RV:** Takes finite or countably infinite distinct values.
  - **Continuous RV:** Can take any value in a continuous interval.

## 3-3 Cumulative Distribution Function (CDF)

- For any RV  $X$ , its CDF  $F_X(x) = P(X \leq x)$ .
- **Properties:**
  - ① Non-decreasing:  $a < b \implies F_X(a) \leq F_X(b)$ .
  - ② Limits:  $\lim_{x \rightarrow -\infty} F_X(x) = 0$ ,  $\lim_{x \rightarrow +\infty} F_X(x) = 1$ .
  - ③ Right-continuous:  $\lim_{h \rightarrow 0^+} F_X(x + h) = F_X(x)$ .
- $P(a < X \leq b) = F_X(b) - F_X(a)$ .

## 3-4 Probability Mass Function (PMF) (Discrete RVs)

- For a discrete RV  $X$  taking values in  $S = \{x_1, x_2, \dots\}$ .
- PMF  $p_X(x) = P(X = x)$ .
- **Properties:**
  - ①  $p_X(x) \geq 0$  for  $x \in S$ , and  $p_X(x) = 0$  if  $x \notin S$ .
  - ②  $\sum_{x \in S} p_X(x) = 1$ .
- Examples: Bernoulli( $p$ ), Binomial( $n, p$ ), Poisson( $\lambda$ ).

## 3-5 Probability Density Function (PDF) (Continuous RVs)

- For a continuous RV  $X$ , its PDF  $f_X(x)$  is such that  $F_X(x) = \int_{-\infty}^x f_X(t)dt$ .
- $f_X(x) = \frac{dF_X(x)}{dx}$  where  $F_X(x)$  is differentiable.
- **Important:**  $f_X(x)$  is NOT  $P(X = x)$ . For continuous RVs,  $P(X = x) = 0$ .
- $P(a < X \leq b) = \int_a^b f_X(x)dx$ .
- **Properties:**
  - 1  $f_X(x) \geq 0$ . (PDF values can be  $> 1$ ).
  - 2  $\int_{-\infty}^{\infty} f_X(x)dx = 1$ .



## 3-6 Expected Value, Variance, Covariance

- **Expected Value (Mean  $E[X]$  or  $\mu_X$ ):**
  - Discrete:  $E[g(X)] = \sum_{x \in S} g(x)p_X(x)$ .
  - Continuous:  $E[g(X)] = \int_{-\infty}^{\infty} g(x)f_X(x)dx$ .
- **Variance ( $Var(X)$  or  $\sigma_X^2$ ):**
  - $Var(X) = E[(X - E[X])^2] = E[X^2] - (E[X])^2$ .
  - Standard Deviation:  $\sigma_X = \sqrt{Var(X)}$ .
- **Covariance ( $Cov(X, Y)$ ):** (Requires joint distribution)
  - $Cov(X, Y) = E[(X - E[X])(Y - E[Y])] = E[XY] - E[X]E[Y]$ .
- **Correlation ( $\rho_{XY}$ ):**  $\frac{Cov(X, Y)}{\sigma_X \sigma_Y}$ .
- PyTorch distributions often have `.mean`, `.variance` properties.

## 3-7 Joint, Marginal, Conditional Distributions

- **Joint Distribution:** Describes simultaneous behavior of multiple RVs.
  - $F_{X,Y}(x,y) = P(X \leq x, Y \leq y)$  (Joint CDF).
  - $p_{X,Y}(x,y) = P(X = x, Y = y)$  (Joint PMF).
  - $f_{X,Y}(x,y)$  (Joint PDF).
- **Marginal Distribution:** Distribution of a subset of RVs, "averaging out" others.
  - $p_X(x) = \sum_y p_{X,Y}(x,y)$ .
  - $f_X(x) = \int f_{X,Y}(x,y) dy$ .
- **Conditional Distribution:** Distribution of one RV given specific values for others.
  - $p_{Y|X}(y|x) = P(Y = y|X = x) = \frac{p_{X,Y}(x,y)}{p_X(x)}$ . (Similar for PDF).
- **Chain Rule:**  $p(x,y) = p(y|x)p(x)$ . Generalizes to many variables, key for graphical models.

## 3-8 Independence and Conditional Independence

- **Independence of RVs ( $X \perp Y$ ):**

- Joint distribution factors:  $p_{X,Y}(x,y) = p_X(x)p_Y(y)$  (or  $f_{X,Y}(x,y) = f_X(x)f_Y(y)$ ).
- Knowing one gives no info about the other:  $p_{Y|X}(y|x) = p_Y(y)$ .
- If independent,  $\text{Cov}(X, Y) = 0$ . (Converse not always true unless jointly Normal).

- **Conditional Independence ( $X \perp Y|Z$ ):**

- $p_{X,Y|Z}(x,y|z) = p_{X|Z}(x|z)p_{Y|Z}(y|z)$ .
- $p_{X|Y,Z}(x|y,z) = p_{X|Z}(x|z)$  (Given  $Z$ ,  $Y$  gives no extra info about  $X$ ).
- Cornerstone of probabilistic graphical models (e.g., Naive Bayes assumes features are conditionally independent given class).

## 3-9 Limit Theorems: LLN and CLT

Sum/average of many i.i.d. RVs tends towards Normal distribution.

### Law of Large Numbers (LLN)

- Let  $X_1, \dots, X_n$  be i.i.d. with mean  $\mu$ . Sample mean  $\bar{X}_n = \frac{1}{n} \sum X_i$ .
- **Weak LLN:**  $\lim_{n \rightarrow \infty} P(|\bar{X}_n - \mu| < \epsilon) = 1$ . (Convergence in probability).
- **Strong LLN:**  $P(\lim_{n \rightarrow \infty} \bar{X}_n = \mu) = 1$ . (Almost sure convergence).
- **Importance:** Justifies Monte Carlo, param estimation, empirical risk.

### Central Limit Theorem (CLT)

- Let  $X_i$  be i.i.d. with mean  $\mu$ , variance  $\sigma^2$ .
- Standardized sum/mean:  $Z_n = \frac{\sum X_i - n\mu}{\sigma\sqrt{n}} = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0, 1)$ .
- So,  $\bar{X}_n \approx N(\mu, \sigma^2/n)$  for large  $n$ .
- **Importance:** Foundation for CIs, explains ubiquity of Normal dist.

# 3-10 Concentration Inequalities

Non-asymptotic bounds on deviation from expectation.

- **Markov's Inequality:** For non-negative RV  $X$ ,  $a > 0$ :

$$\mathbb{P}(X \geq a) \leq \frac{\mathbb{E}[X]}{a}$$

- General but often loose.  $X$  must be non-negative.
- **Chebyshev's Inequality:** For RV  $X$  with mean  $\mu$ , variance  $\sigma^2$ ,  $k > 0$ :

$$\mathbb{P}(|X - \mu| \geq k\sigma) \leq \frac{1}{k^2}$$

- Stronger than Markov, uses variance. Applies to any distribution with finite mean/var.

# 3-10 Concentration Inequalities

Non-asymptotic bounds on deviation from expectation.

- **Extended Markov's Inequality:**

- If  $\varphi$  is a nondecreasing nonnegative function,  $X$  is a (not necessarily nonnegative) random variable, and  $\varphi(a) > 0$ , then

$$\mathbb{P}(X \geq a) \leq \frac{\mathbb{E}[\varphi(X)]}{\varphi(a)}.$$

- **Corollary (Higher Moments):** For any RV  $X$ , integer  $n \geq 1$ , and  $a > 0$ :

$$\mathbb{P}(|X| \geq a) \leq \frac{\mathbb{E}[|X|^n]}{a^n}.$$

(This follows by applying the extended form to the RV  $|X|$  with  $\varphi(x) = x^n$  for  $x \geq 0$ ).

## 3-10 Concentration Inequalities

Further bounds and extensions.

- **Hoeffding's Inequality:** For sum  $S_n$  of  $n$  independent *bounded* RVs  $X_i \in [a_i, b_i]$ :

$$\mathbb{P}(S_n - \mathbb{E}[S_n] \geq t) \leq \exp\left(-\frac{2t^2}{\sum (b_i - a_i)^2}\right)$$

For i.i.d. Bernoulli( $p$ ) variables,  $\bar{X}_n = \hat{p}$ :

$$\mathbb{P}(|\hat{p} - p| \geq \epsilon) \leq 2 \exp(-2n\epsilon^2).$$

- Tighter for bounded RVs, exponential decay. **Crucial for generalization bounds in ML.**
- **Others:** Chernoff, Bernstein, McDiarmid (more specialized/tighter).

# Mathematical Statistics

Learning from Data in the Presence of Uncertainty

## Core Idea

Using data to infer properties of an unknown underlying probability distribution or model.

### Probability Theory:

- Model  $\rightarrow$  Data
- Known model, predict outcomes.
- Deductive.
- "If coin is fair,  $P(7H \text{ in } 10)$ ?"

### Statistics:

- Data  $\rightarrow$  Model
- Unknown model, infer from data.
- Inductive.
- "Given  $7H$  in 10, is coin fair?"

## Why it's important for ML/DL?

- Parameter Estimation, Model Building/Selection, Model Evaluation, Uncertainty Quantification, Understanding Generalization.



## 4-1 Statistical Inference: The Goal

Deduce properties of an underlying probability distribution from data.

- **Point Estimation:** Single "best" value for an unknown parameter (e.g., mean  $\mu$ ).
- **Interval Estimation:** Range of plausible values (e.g., confidence interval for  $\mu$ ).
- **Hypothesis Testing:** Decide between competing claims (e.g., is  $\mu = 0$  vs  $\mu \neq 0$ ?).
- **Prediction:** Predict future observations.

*We primarily focus on point estimation for ML model training.*

## 4-2 Samples, Statistics, Estimators

- **Random Sample:**  $X_1, \dots, X_n$  i.i.d. from  $f(x|\theta)$  or  $p(x|\theta)$ .
  - Observed data  $x_1, \dots, x_n$  are realizations.
- **Statistic:** Function  $T(X_1, \dots, X_n)$  of sample, not depending on unknown  $\theta$ .
  - E.g., sample mean  $\bar{X}$ , sample variance  $S^2$ .
- **Estimator  $\hat{\theta}$ :** Statistic used to estimate  $\theta$ . It's an RV.
- **Estimate:** Specific value  $\hat{\theta}(x_1, \dots, x_n)$  from data.

### Desirable Properties of Estimators

- **Unbiasedness:**  $E[\hat{\theta}] = \theta$ . (Bias:  $E[\hat{\theta}] - \theta$ ).
- **Efficiency:** Smaller variance among unbiased estimators.
- **Consistency:**  $\hat{\theta}_n \xrightarrow{P} \theta$  as  $n \rightarrow \infty$ .
- **Sufficiency:** Captures all info about  $\theta$  from sample.

## 4-3 Point Estimation: Maximum Likelihood Estimation (MLE)

Find parameter value that makes observed data most probable.

- **Likelihood Function**  $L(\theta|\mathbf{x})$ : Joint PMF/PDF of sample, viewed as function of  $\theta$  for fixed data  $\mathbf{x}$ .

$$L(\theta|\mathbf{x}) = \prod_{i=1}^n f(x_i|\theta) \quad (\text{or } p(x_i|\theta))$$

- $L(\theta|\mathbf{x})$  is NOT a probability dist. for  $\theta$ .
- **MLE Principle:**  $\hat{\theta}_{MLE} = \arg \max_{\theta} L(\theta|\mathbf{x})$ .
- **Log-Likelihood**  $\ell(\theta|\mathbf{x})$ :  $\log L(\theta|\mathbf{x}) = \sum_{i=1}^n \log f(x_i|\theta)$ .
  - Easier to work with (products to sums, numerical stability).
  - $\hat{\theta}_{MLE} = \arg \max_{\theta} \ell(\theta|\mathbf{x})$ .
- **Finding MLE:** Differentiate  $\ell(\theta|\mathbf{x})$  w.r.t  $\theta$ , set to 0, solve. (Or numerical opt.)

## 4-4 MLE Examples

### Example 1: Bernoulli parameter $p$

Data  $X_1, \dots, X_n \sim \text{Bernoulli}(p)$ .  $k = \sum x_i$ .

$\ell(p|\mathbf{x}) = k \log p + (n - k) \log(1 - p)$ .  $\implies \hat{p}_{MLE} = k/n = \bar{x}$  (sample proportion).

### Example 2: Mean $\mu$ of Normal $N(\mu, \sigma^2)$ ( $\sigma^2$ known)

$\ell(\mu|\mathbf{x}, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum (x_i - \mu)^2$ .  $\implies \hat{\mu}_{MLE} = \frac{1}{n} \sum x_i = \bar{x}$  (sample mean).

## 4-5 MLE Properties

### Properties of MLEs (under regularity conditions)

Consistent, Asymptotically Normal, Asymptotically Efficient, Invariance.

- May be biased for finite samples (e.g.,  $\hat{\sigma}_{MLE}^2$  for Normal).

### MLE in Machine Learning

Many loss functions are Negative Log-Likelihoods (NLL). Minimizing NLL  $\equiv$  MLE.

- MSE Loss (LinReg with Gaussian noise)  $\equiv$  MLE.
- Cross-Entropy Loss (LogReg/Classification)  $\equiv$  MLE.

# 4-6 Bayesian Inference & Maximum A Posteriori (MAP)

## Incorporating Prior Beliefs

### Frequentist (MLE):

- $\theta$  is fixed, unknown.
- Data is random.

### Bayesian (MAP):

- $\theta$  is a random variable with a distribution.
- Data is observed, fixed.

## Bayes' Theorem for Parameters

$$\text{Posterior } P(\theta|\mathbf{x}) = \frac{\text{Likelihood } P(\mathbf{x}|\theta) \times \text{Prior } P(\theta)}{\text{Evidence } P(\mathbf{x})}$$

- **Prior**  $P(\theta)$ : Beliefs about  $\theta$  *before* data.
- **Posterior**  $P(\theta|\mathbf{x})$ : Updated beliefs about  $\theta$  *after* data.
- **MAP Estimation**: Point estimate maximizing posterior.

$$\begin{aligned}\hat{\theta}_{MAP} &= \arg \max_{\theta} P(\theta|\mathbf{x}) = \arg \max_{\theta} [P(\mathbf{x}|\theta)P(\theta)] \\ &= \arg \max_{\theta} [\ell(\theta|\mathbf{x}) + \log P(\theta)]\end{aligned}$$

## 4-7 MAP, MLE, and Regularization

- **MAP vs MLE:**

- If prior  $P(\theta)$  is uniform (flat),  $\log P(\theta)$  is constant  $\implies \hat{\theta}_{MAP} = \hat{\theta}_{MLE}$ .
- As data  $n \rightarrow \infty$ , likelihood dominates prior  $\implies \hat{\theta}_{MAP} \rightarrow \hat{\theta}_{MLE}$ .

- **MAP and Regularization in ML:**

- Minimizing (NLL + Regularizer)  $\equiv$  MAP estimation.
- **L2 Regularization (Ridge/Weight Decay):**

$$\text{NLL}(\mathbf{w}|\mathbf{x}) + \lambda \|\mathbf{w}\|_2^2$$

Equivalent to MAP with Gaussian prior  $N(0, \sigma_p^2)$  on weights  $\mathbf{w}$ .  
( $\lambda \propto 1/\sigma_p^2$ )

- **L1 Regularization (Lasso):**

$$\text{NLL}(\mathbf{w}|\mathbf{x}) + \lambda \|\mathbf{w}\|_1$$

Equivalent to MAP with Laplacian prior on weights  $\mathbf{w}$ . Promotes sparsity.

- Conceptual PyTorch MAP: Add `log_prior` term to loss: `loss = -(log_likelihood + log_prior)`.

## 4-8 Hypothesis Testing & Confidence Intervals (Briefly)

- **Hypothesis Testing:** Formal procedure to decide between two competing statements ( $H_0$ : null,  $H_A$ : alternative) about a population based on sample data.
  - Calculate test statistic  $\rightarrow$  p-value.
  - **p-value:**  $P(\text{observed data or more extreme} \mid H_0 \text{ is true})$ .
  - If p-value  $\leq$  significance level  $\alpha$  (e.g., 0.05), reject  $H_0$ .
  - ML: Model comparison, feature significance.
- **Confidence Intervals (CI):** Range of values likely to contain true parameter value with certain confidence (e.g., 95% CI).
  - If repeated sampling, 95% of CIs would contain true parameter.
  - Measures uncertainty around point estimate.
- *Crucial for rigorous model evaluation, A/B testing, scientific reporting in ML.*



## 4-9 Bias-Variance Tradeoff (Statistical Perspective)

### Decomposing Expected Prediction Error

Assume  $Y = f_{true}(\mathbf{X}) + \epsilon$ , with  $E[\epsilon] = 0$ ,  $Var(\epsilon) = \sigma_\epsilon^2$ . Our learned model is  $\hat{f}(\mathbf{X})$ . Expected Prediction Error (EPE) at  $\mathbf{x}_0$  for squared error loss:

$$EPE(\mathbf{x}_0) = \underbrace{(E_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(\mathbf{x}_0)] - f_{true}(\mathbf{x}_0))^2}_{\text{Bias}^2(\hat{f}(\mathbf{x}_0))} + \underbrace{E_{\mathcal{D}}[(\hat{f}_{\mathcal{D}}(\mathbf{x}_0) - E_{\mathcal{D}}[\hat{f}_{\mathcal{D}}(\mathbf{x}_0)])^2]}_{\text{Variance}(\hat{f}(\mathbf{x}_0))} + \sigma_\epsilon^2$$

The final term is the Irreducible Error  $\sigma_\epsilon^2$ .

- **Bias:** Systematic error. Average difference between model's prediction and true function. (Underfitting: high bias).
- **Variance:** Model's sensitivity to specific training data. Variability of predictions if retrained on different datasets. (Overfitting: high variance).
- **Irreducible Error:** Noise inherent in data. Cannot be reduced.

## 4-9 Bias-Variance Tradeoff (Statistical Perspective)

### The Tradeoff

- Increasing model complexity  $\implies \downarrow$  Bias,  $\uparrow$  Variance.
- Decreasing model complexity  $\implies \uparrow$  Bias,  $\downarrow$  Variance.
- Goal: Find balance to minimize total error. (Cross-validation, regularization, ensembles).

Thank You!